

Numerical stability and fast-math

Speeding up LHCb software through compilation optimization

Oscar Buon

July 6, 2023

Supervisor: Sebastien Ponce

Table of contents

- 1 Introduction
 - Floating-point arithmetic
 - IEEE 754
- 2 Correctness of some computes
 - Floats equality
 - Exception cases
- 3 Fast-math
 - Exception handling
 - Approximation
- 4 Results
 - Improvement
 - Difference

- 1 Introduction
 - Floating-point arithmetic
 - IEEE 754
- 2 Correctness of some computes
 - Floats equality
 - Exception cases
- 3 Fast-math
 - Exception handling
 - Approximation
- 4 Results
 - Improvement
 - Difference

Floating-point arithmetic

$$\begin{aligned}
 (0.75)_{10} &= (+7.5 \times 10^{-1})_{10} \\
 &= \underbrace{(+)}_{\text{sign}} \underbrace{1.1}_{\text{mantissa}} \times 2^{\underbrace{-1}_{\text{exponent}}} = (0.11)_2
 \end{aligned}$$

IEEE 754

- 32 bits and 64 bits formats
- Subnormal numbers
- Special values: $-\text{Inf}$, $+\text{Inf}$, NaN
- Rounding rules
- Exception handling
- We should always consider floats as approximations !

- 1 Introduction
 - Floating-point arithmetic
 - IEEE 754
- 2 Correctness of some computes
 - Floats equality
 - Exception cases
- 3 Fast-math
 - Exception handling
 - Approximation
- 4 Results
 - Improvement
 - Difference

Floats equality

- Never check for float equality.
- GCC flag `-Wfloat-equal`. Ex. in LHCb:
 - `0 == ip`
 - `0.5 == ip`
 - `lhs.m_energy == rhs.m_energy`

Exception cases

```
if (f(x) != 0.)  
    return 5./f(x)
```

- Maybe the two $f(x)$ will not be the same.


```
if ((float y = f(x)) != 0.)  
    return 5./y
```

- Maybe $f(x)$ will still be calculated twice.
- if $f(x)$ is close to 0 then $5./y$ could be $+\text{Inf}$.

We should use:

- Using `isfinite` after the compute.

There is also:

- The trapping system (registering a handler with `signal()`).
- The signaling system via
`std::numeric_limits<T>::signaling_NaN`.

- 1 Introduction
 - Floating-point arithmetic
 - IEEE 754
- 2 Correctness of some computes
 - Floats equality
 - Exception cases
- 3 **Fast-math**
 - Exception handling
 - Approximation
- 4 Results
 - Improvement
 - Difference

Principle

- Making mathematically valid optimizations but not respecting Standard.
- Different results but not more wrong than without `fast-math`.

no-math-errno

- Do not set `errno` after calling math functions that are executed with a single instruction, e.g., `sqrt`.

no-signaling-nans

- Compile code assuming that IEEE signaling NaNs may not generate user-visible traps during floating-point operations.
- Enables optimizations that may change the number of exceptions visible with signaling NaNs.
- Enabled by default.

no-trapping-math

- Compile code assuming that floating-point operations cannot generate user-visible traps. These traps include division by zero, overflow, underflow, inexact result and invalid operation.
- Implies no-signaling-nans.

finite-math-only

- Allow optimizations for floating-point arithmetic that assume that arguments and results are not NaNs or $\pm\text{Infs}$.
- Compiler can replace `isnan` by `false`. Actually it is an undefined behavior.
- Main source of bugs from `fast-math`.

no-rounding-math

- Round-to-zero for all floating point to integer conversions, and round-to-nearest for all other arithmetic truncations.
- Enabled by default.

no-signed-zeros

- Allow optimizations for floating-point arithmetic that ignore the signedness of zero.
- Ex. $0.0+x$ or $0.0*x$.

associative-math

- Allow re-association of operands in series of floating-point operations.
- Can optimize $2.0*x*3.0$ in $6.0*x \implies$ make some computes at compile time.
- May allow a better vectorization and use of *FMA*.
- Needs no-signed-zeros and no-trapping-math.

reciprocal-math

- Allow the reciprocal of a value to be used instead of dividing by the value if this enables optimizations.
- Ex. Replacing x/y by $x*(1/y)$.
- Decreases precision.

unsafe-math-optimizations

- Allow optimizations for floating-point arithmetic that (a) assume that arguments and results are valid and (b) may violate IEEE or ANSI standards.
- When used at link time, it may include libraries or startup files that change the default FPU control word or other similar optimizations.
- Can affect dynamically included libraries.
- Enables `-fno-signed-zeros`, `-fno-trapping-math`, `-fassociative-math` and `-freciprocal-math`.

- 1 Introduction
 - Floating-point arithmetic
 - IEEE 754
- 2 Correctness of some computes
 - Floats equality
 - Exception cases
- 3 Fast-math
 - Exception handling
 - Approximation
- 4 Results
 - Improvement
 - Difference

Improvement

Test: hlt2_pp_thor (2×1000 events)

Optimization	Improvement	Confidence interval (2σ)
Fast-math ¹	0%	$\pm 0\%$
Associative-math	4.73%	$\pm 0.87\%$
Fast-math + LTO & PGO	12.04%	$\pm 0.73\%$

¹without finite-math-only

FP Arithmetic [®] :	10.9%
FP x87 [®] :	0.0%
FP Scalar [®] :	6.7%
FP Vector [®] :	4.2%


Figure: Reference

FP Arithmetic [®] :	9.5%
FP x87 [®] :	0.0%
FP Scalar [®] :	5.9%
FP Vector [®] :	3.6%

Figure: Associative-math

Difference

Difference with the reference (all counters):



Conclusion

- Enabling `-Wfloat-equal`.
- Enabling `fast-math` at least for checking instability.
- Other advantages:
 - Better precision in most cases (*FMA* and `double-precision-constant`).
 - Coding clear equations and letting compiler optimize them.
 - GPU already using similar principles.
 - Make it clear that floats are approximations.