

# Numerical stability and fast-math

## Speeding up LHCb software through compilation optimization

Oscar Buon

July 7, 2023

Supervisor: Sebastien Ponce

# Table of contents

- 1 Introduction
  - Floating-point arithmetic
  - IEEE 754
- 2 Correctness of some computations
  - Floats equality
  - Exception cases
- 3 Fast-math
  - Exception handling
  - Approximation
- 4 Results
  - Improvement
  - Differences in the counters

- 1 Introduction
  - Floating-point arithmetic
  - IEEE 754
- 2 Correctness of some computations
  - Floats equality
  - Exception cases
- 3 Fast-math
  - Exception handling
  - Approximation
- 4 Results
  - Improvement
  - Differences in the counters

# Floating-point arithmetic

$$\begin{aligned}
 (0.75)_{10} &= (+7.5 \times 10^{-1})_{10} \\
 &= (+11 \times 2^{-2})_2 \\
 &= (\underbrace{+}_{\text{sign}} \underbrace{1.1}_{\text{mantissa}} \times 2^{\overbrace{-1}^{\text{exponent}}})_2 = (0.11)_2
 \end{aligned}$$

# Non-nominal case

$$(0.8)_{10} \approx (1.1001100110011001101 \times 2^{-1})_2$$

The representation is an approximation, because the number of bits is limited.

# IEEE 754

- 32 bits ( $\sim 7.2$  digits) and 64 bits ( $\sim 15.9$  digits) formats ;
- Special rules:
  - Subnormal numbers (numbers very close to 0) ;
  - Special values:  $-\text{Inf}$ ,  $+\text{Inf}$ , NaN ;
  - Rounding rules ;
  - Exception handling ;
- We should always consider floats as approximations !

- 1 Introduction
  - Floating-point arithmetic
  - IEEE 754
- 2 Correctness of some computations
  - Floats equality
  - Exception cases
- 3 Fast-math
  - Exception handling
  - Approximation
- 4 Results
  - Improvement
  - Differences in the counters

# Floats equality

- Never check for float equality.
- GCC flag `-Wfloat-equal`. Ex. in LHCb:
  - `0 == ip`
  - `0.5 == ip`
  - `lhs.m_energy == rhs.m_energy`



# Exception cases

```
if (f(x) != 0.)  
    return 5./f(x);  
else  
    ...
```

- Maybe the two  $f(x)$  will not be the same.
- Due to different optimizations.

```
float y = f(x);  
if (y != 0.)  
    return 5./y;  
else  
    ...
```

- Maybe  $f(x)$  will still be calculated twice.
- If  $f(x)$  is close to 0 then  $5./y$  could be  $+\text{Inf}$ .

A good way is using `isfinite` after the computation.

```
float result = 5./f(x);  
if (std::isfinite(result))  
    return result;  
else  
    ...
```

There are also:

- The trapping system (registering a handler with C function `signal(handler)`);
- The signaling system via  
`std::numeric_limits<T>::signaling_NaN`;

but they are old methods and are not recommended.

- 1 Introduction
  - Floating-point arithmetic
  - IEEE 754
- 2 Correctness of some computations
  - Floats equality
  - Exception cases
- 3 **Fast-math**
  - Exception handling
  - Approximation
- 4 Results
  - Improvement
  - Differences in the counters

# Principle

- Set of flags.
- Making mathematically valid optimizations but not respecting Standard:
  - $a + b - a = b$ .
  - With  $a = 10^6$  and  $b = 10^{-6}$ ,  
Standard gives 0, fast-math gives  $10^{-6}$ .
- Different results but not more wrong than without fast-math.

<https://godbolt.org/z/jqz8P8vjj>

## no-math-errno

- Do not set `errno` after calling math functions that are executed with a single instruction, e.g., `sqrt`.

## no-signaling-nans

- Compile code assuming that IEEE signaling NaNs may not generate user-visible traps during floating-point operations.
- Enables optimizations that may change the number of exceptions visible with signaling NaNs.
- Enabled by default.



# no-trapping-math

- Compile code assuming that floating-point operations cannot generate user-visible traps. These traps include division by zero, overflow, underflow, inexact result and invalid operation.
- Implies no-signaling-nans.

# finite-math-only

- Allow optimizations for floating-point arithmetic that assume that arguments and results are not NaNs or  $\pm\text{Infs}$ .
- Compiler can replace `isnan` by `false`. Actually it is an undefined behavior.
- Main source of issues from `fast-math`.

# no-signed-zeros

- Allow optimizations for floating-point arithmetic that ignore the signedness of zero.
- Ex.  $0.0+x$  or  $0.0*x$ .

## associative-math

- Allow re-association of operands in series of floating-point operations.
- Can optimize  $2.0 * x * 3.0$  in  $6.0 * x \implies$  make some computations at compile time.
- May allow a better vectorization and use of *FMA*.
- Needs no-signed-zeros and no-trapping-math.

# reciprocal-math

- Allow the reciprocal of a value to be used instead of dividing by the value if this enables optimizations.
- Ex. Replacing  $x/y$  by  $x*(1/y)$ .
- May decrease precision.

# unsafe-math-optimizations

- Allow optimizations for floating-point arithmetic that
  - ① assume that arguments and results are valid and
  - ② may violate IEEE or ANSI standards.
- When used at link time, it may include libraries or startup files that change the default FPU control word or other similar optimizations.
- Can affect dynamically included libraries.
- Enables `-fno-signed-zeros`, `-fno-trapping-math`, `-fassociative-math` and `-freciprocal-math`.

- 1 Introduction
  - Floating-point arithmetic
  - IEEE 754
- 2 Correctness of some computations
  - Floats equality
  - Exception cases
- 3 Fast-math
  - Exception handling
  - Approximation
- 4 Results
  - Improvement
  - Differences in the counters

# Improvement

Test: hlt2\_pp\_thor

Optimization	Improvement	Confidence interval ( $2\sigma$ )
Fast-math <sup>1</sup>	5.06%	$\pm 0.98\%$
Associative-math only	4.73%	$\pm 1.51\%$
Fast-math <sup>1</sup> + LTO & PGO	11.02%	$\pm 0.98\%$

---

<sup>1</sup>without finite-math-only and unsafe-math-optimizations



Time used by CPU for computing floats:

FP Arithmetic <sup>®</sup> :	10.9%
FP x87 <sup>®</sup> :	0.0%
FP Scalar <sup>®</sup> :	6.7%
FP Vector <sup>®</sup> :	4.2%

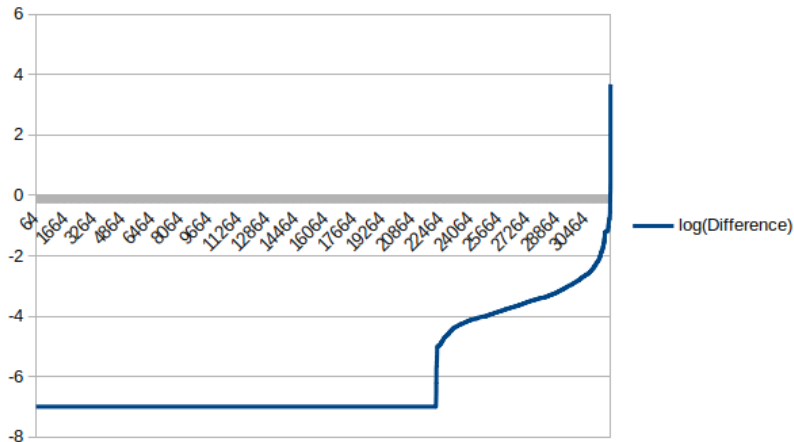
Figure: Reference

FP Arithmetic	9.5%
FP x87	0.0%
FP Scalar	5.9%
FP Vector	3.6%

Figure: Associative-math

# Differences in the counters

Difference with the reference (all counters):



# In a first time

- Enabling `-Wfloat-equal`.
- Enabling `fast-math` in one slot for instability checking.
  - Forget about `finite-math-only` and `unsafe-math-optimizations`.

# In a second time

- Switch to fast-math for production.
  - Sometimes better precision (*FMA* and *double-precision-constant*).
  - Coding clear equations and letting compiler optimize them.
  - GPU already using similar principles.
  - Dropping features that are legacies.
  - Assuming floats are approximations.