



Profile-guided optimization

Speeding up LHCb software through compilation optimization

Oscar Buon

June 19, 2023

Supervisor: Sébastien Ponce

CERN

Intergovernmental particle physics on France-Switzerland border.
About 15,000 people working at CERN.
Made discoveries that led to Nobel Prizes.
World Wide Web was invented at CERN.



Large Hadron Collider (LHC)

- Large Hadron Collider: particle collider
- 27km (biggest in the world)
- $\sim 100m$ underground



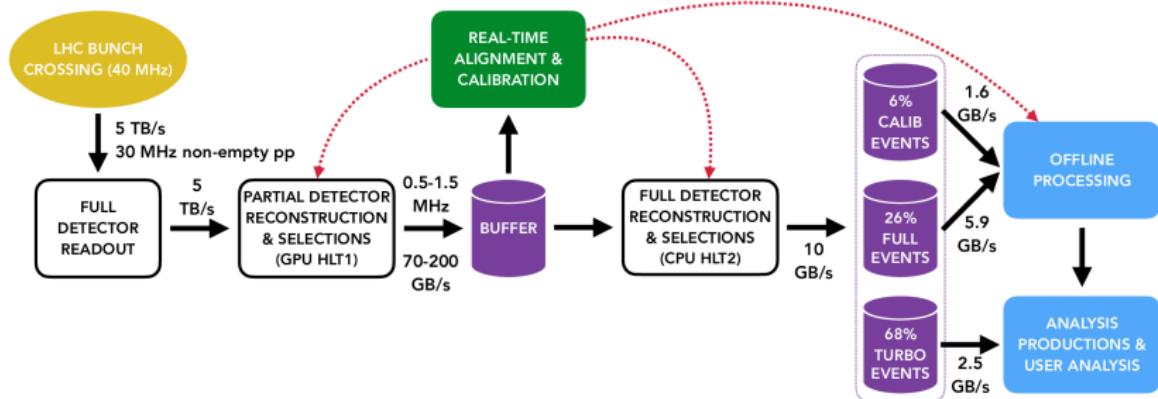
LHCb

- One of the 4 main experiments installed on the LHC.
- More than 1,200 people working for the collaboration.
- Studies asymmetry between matter and antimatter via b-physics. Collision of hadrons (heavy particles).

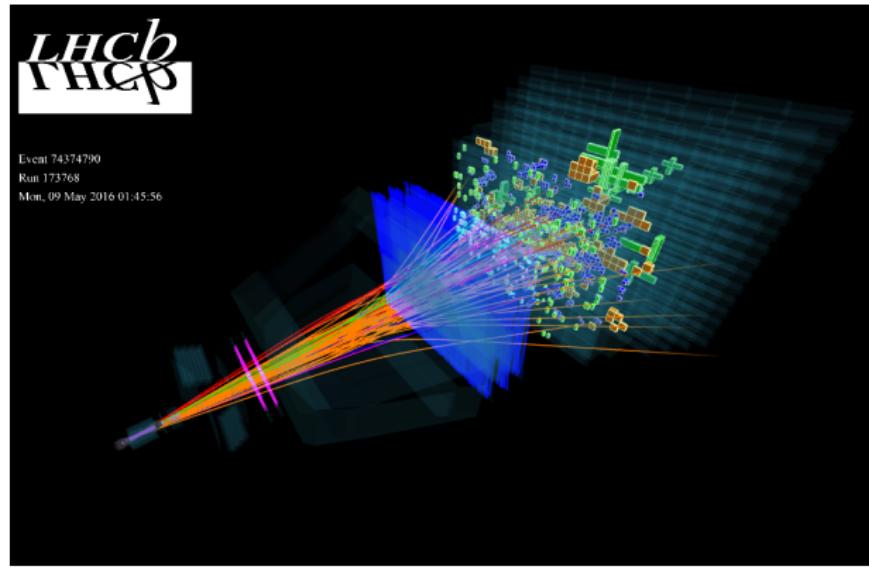


Physics Computing

Important computing infrastructure.



Computing group



The computing group has to maintain this infrastructure for LHCb.

Software

The software is a stack of programs:

- Common base shared with other experiments.
- LHCb specific programs.
- Reconstruction, simulation...

Millions lines of code which some are 30 years old. Code mainly written by non-software engineers.

All of this is running on thousands of multithreaded Linux servers.

Table of contents

1 Graphviz for printing graphs

2 Static libraries and modules

- Library types
- Static libraries
- Static modules

3 Profile-guided optimization

- Profile-guided optimization
- Link-time optimization
- Final building pipeline

1 Graphviz for printing graphs

2 Static libraries and modules

- Library types
- Static libraries
- Static modules

3 Profile-guided optimization

- Profile-guided optimization
- Link-time optimization
- Final building pipeline

Graphviz for printing graphs

Printing graphs was useful to have an idea of the dependencies.
Graphviz in CMake to create dependency graphs:

- CMake arg: `--graphviz=path/to/files.dot`
- Convert to SVG:
`dot -Tsvg -o path/to/file.svg path/to/file.dot`
- Make arg: `CMAKEFLAGS="--graphviz=path/to/files.dot"`
- Setting `GRAPHVIZ_EXTERNAL_LIBS` to FALSE is useful.



1 Graphviz for printing graphs

2 Static libraries and modules

- Library types
- Static libraries
- Static modules

3 Profile-guided optimization

- Profile-guided optimization
- Link-time optimization
- Final building pipeline

Library types

- STATIC: archive of object files that are merged to the executable at link time.
- SHARED: .so (Linux) file that is automatically linked to the executable at runtime (when starting it). Allows several executables to link with a library installed on the system.
- MODULE: same as shared but the module is loaded on demand via the `dlopen` function, so not automatically.

Static libraries

Gaudi provides CMake functions to create libraries, modules or executables. Creating a new function for static libraries by copying the one for shared libraries:

`gaudi_add_library` ⇒ `gaudi_add_static_library`

And inside replacing SHARED by STATIC in add_library function call.

Static modules

By using previous static libraries. Need to add:

- `-Wl,--whole-archive` because without the linker removes what seems to be unused code.
- `-Wl,--allow-multiple-definition` because `-Wl,--whole-archive` causes symbols to be included several times.
- `-Wl,--export-dynamic` to allow functors to access to symbols.

These flags seem to be Hacks and are not recommended to be used. There may still be bugs, especially with functors.

The final executable goes from $\sim 20kB$ to $2.5GB$ (opt+g).
Standard method to run test with python not working. Need to run directly the executable with a json options file:

```
build.x86_64_v3-centos7-gcc11+detdesc-opt+g/run
./build.x86_64_v3-centos7-gcc11+detdesc-opt+
g/Gaudi/Gaudi/Gaudi_static_options.json
```

1 Graphviz for printing graphs

2 Static libraries and modules

- Library types
- Static libraries
- Static modules

3 Profile-guided optimization

- Profile-guided optimization
- Link-time optimization
- Final building pipeline

Profile-guided optimization

The compiler uses some heuristics to optimize some elements:

- Inlining
- Block ordering
- Register allocation
- Virtual call speculation
- Dead code separation

Makes programs running several times faster. But a few times these heuristics are wrong.

A better way for the compiler should be knowing running behavior.
The principle of PGO is:

- Compiling the program with instrumentation.
- Running it to create profiles (counters).
- Recompiling the program with the profiles.

Link-time optimization

C programm composed of translation units (1 translation unit \approx 1 .c file).

Compiler cannot optimize accross them by default.

LTO allows the linker to perform optimizations that take account of all translation units.

Final building pipeline

- Compiling the program with `-fprofile-generate`
- Running it to create profiles
- Recompiling the program with
`-flto -fprofile-use -fprofile-correction`

Conclusion

Test: hlt2_pp_thor (6×1000 events)

Optimization	Acceleration	Confidence interval (2σ)
LTO	0.17%	$\pm 1.12\%$
LTO & PGO	6.74%	$\pm 1.44\%$
Static LTO	0.87%	$\pm 0.60\%$
Static LTO & PGO	6.88%	$\pm 0.83\%$

- Using LTO & PGO makes the programm running faster.
- Using static libraries and modules doesn't seem to be usefull and leads to some bugs.

Context

Worked on a stack with one main CMake:

<https://gitlab.cern.ch/clemenci/lhcb-super-project-template/>

Target: x86_64_v3-centos7-gcc11+detdesc-opt+g